

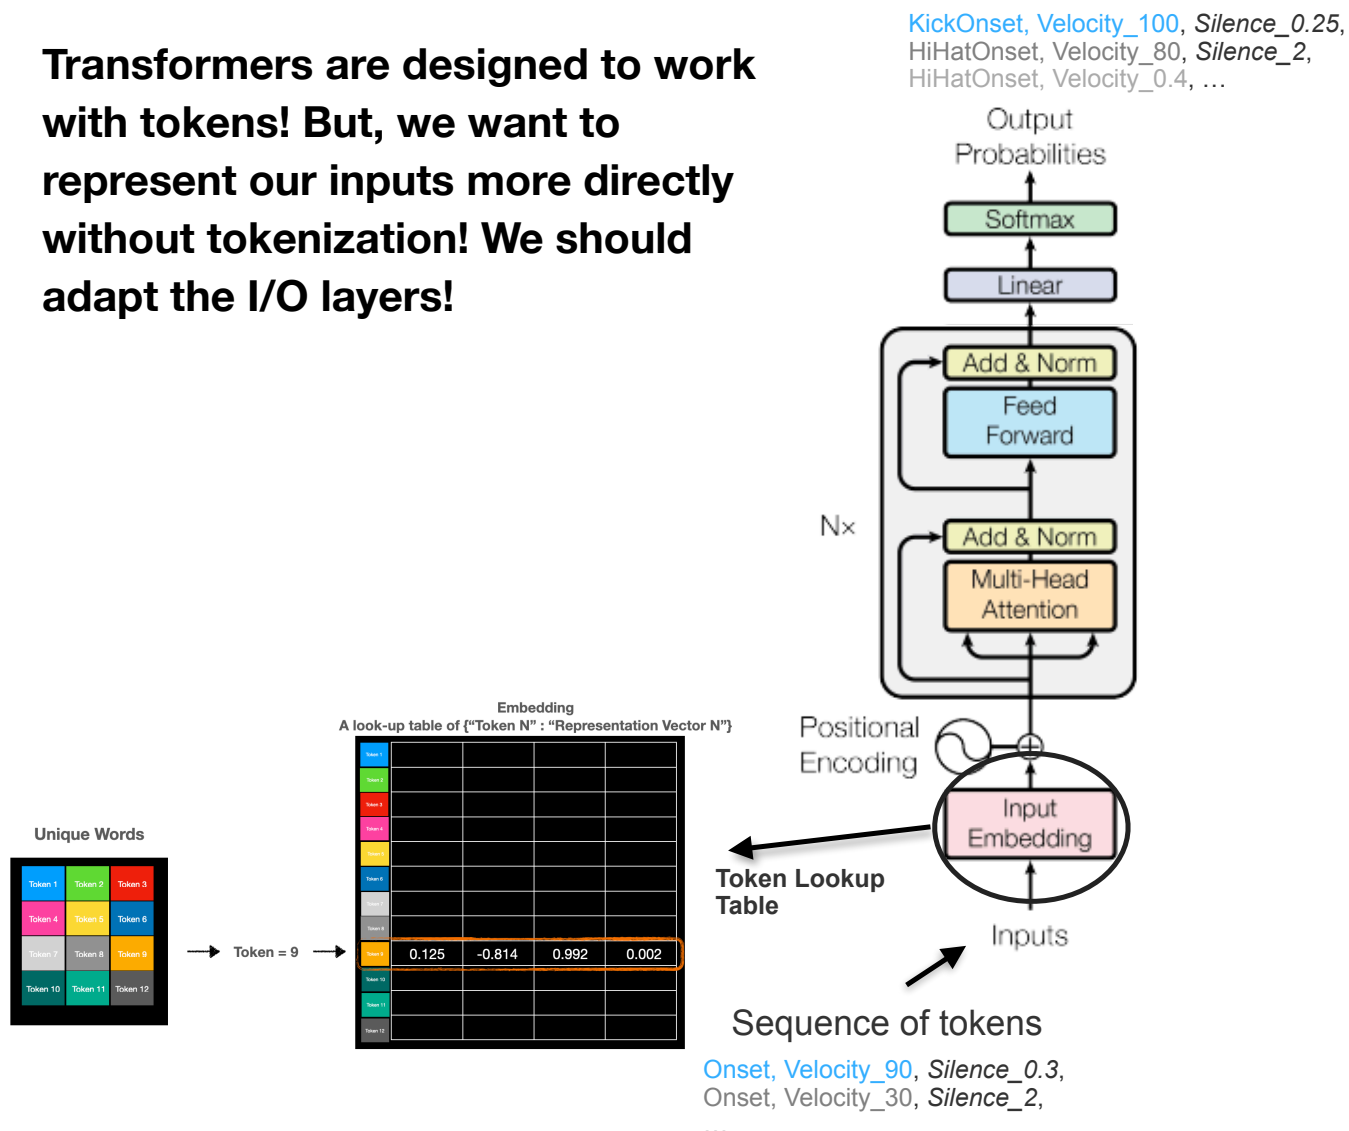
Designing NN-Based Generative Models of Music (Part 3)

Computational Music Creativity, SMC (2022/23)

Prepared by Behzad Haki, February 2023

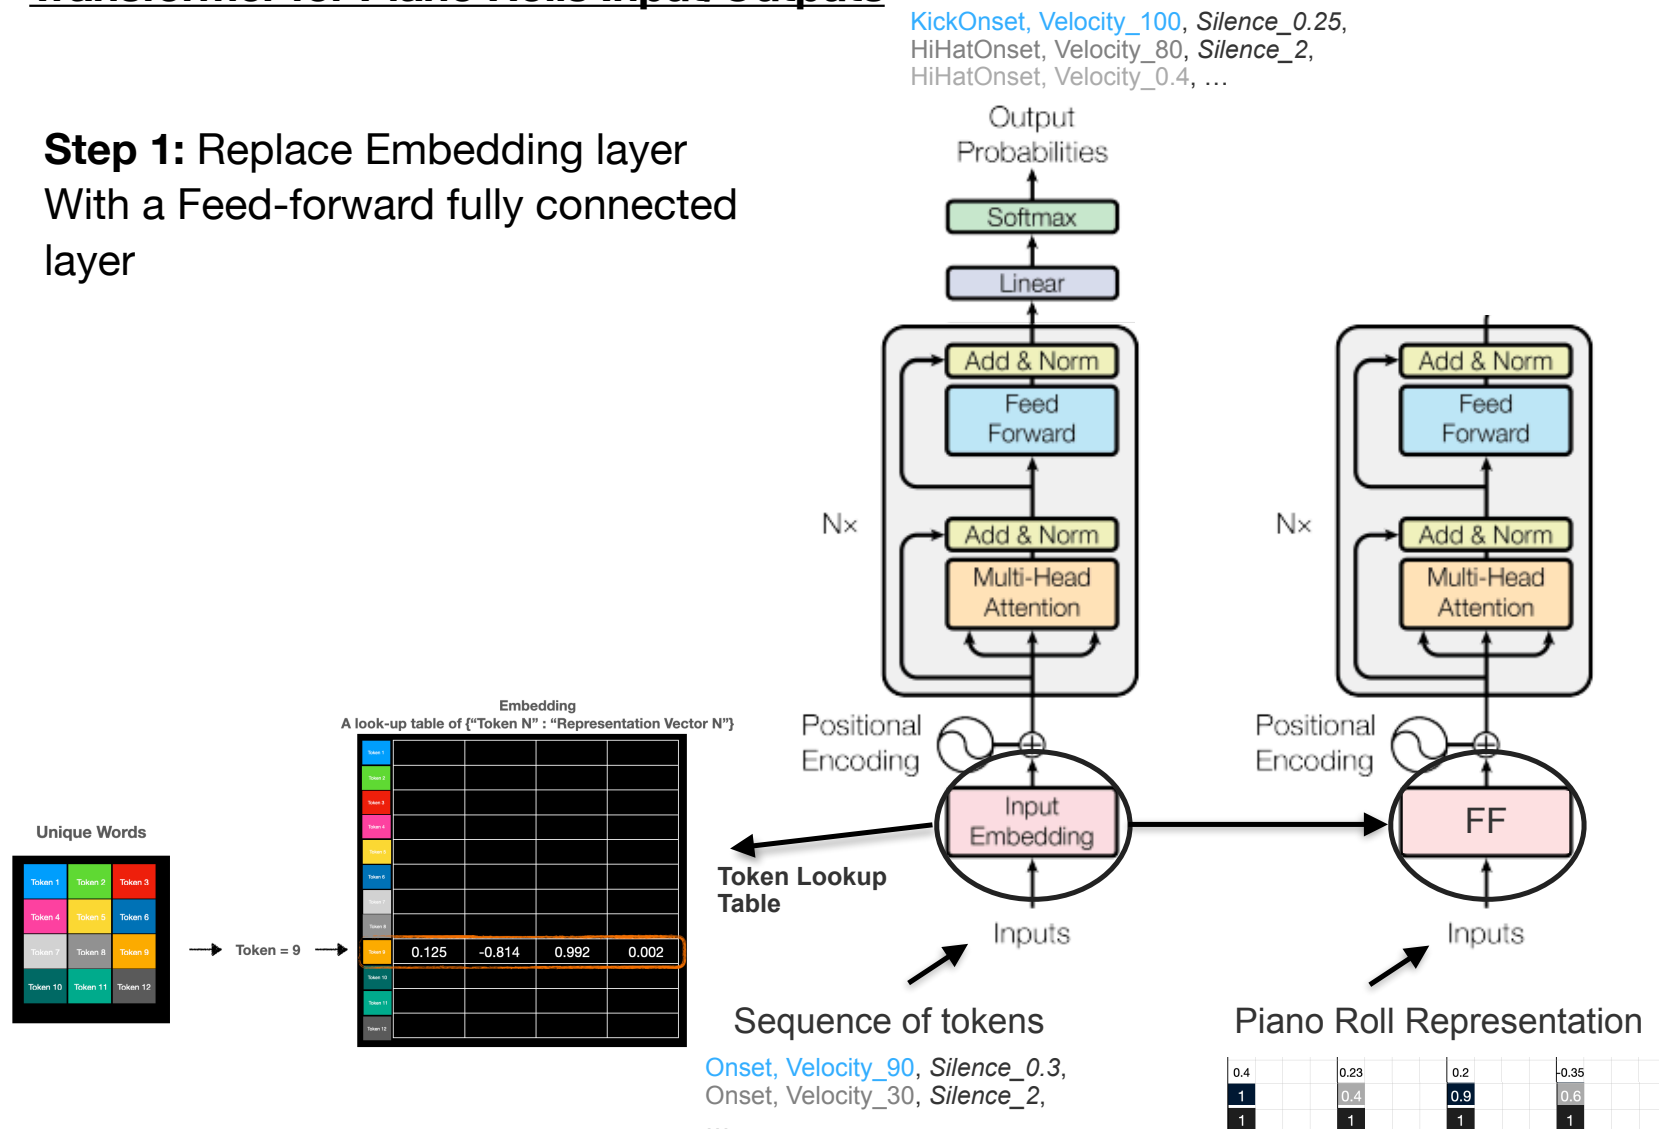
Architecture

Transformers are designed to work with tokens! But, we want to represent our inputs more directly without tokenization! We should adapt the I/O layers!



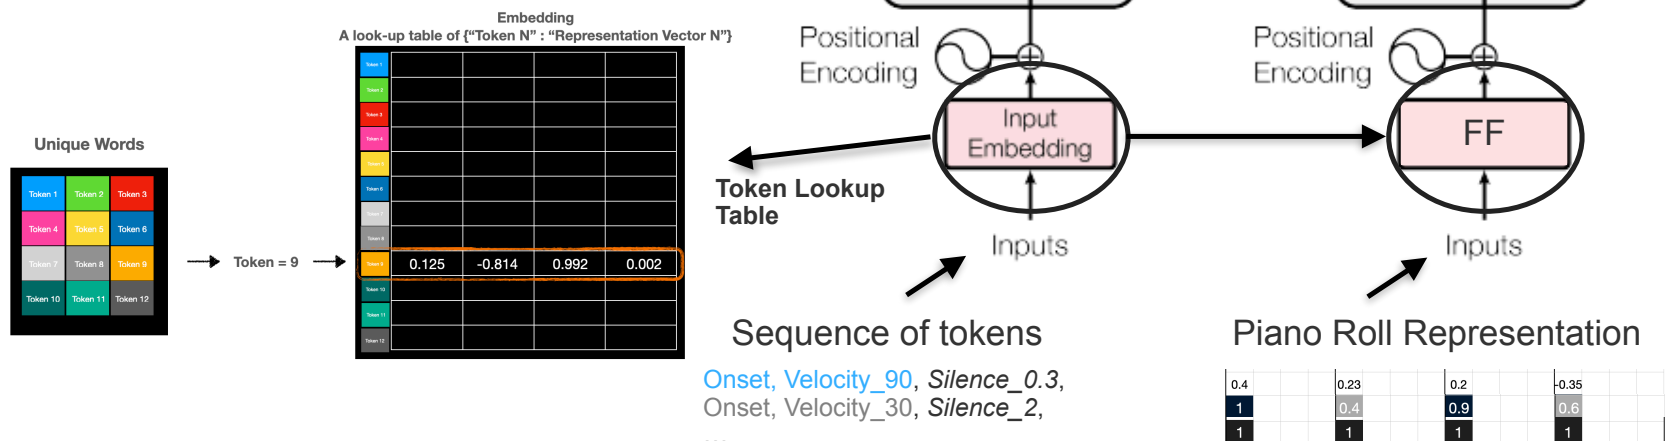
Architecture: Transformer for Piano Rolls Input/Outputs

Step 1: Replace Embedding layer
With a Feed-forward fully connected
layer



Architecture: Transformer for Piano Rolls Input/Outputs

Step 2: Modify the output layer



Architecture

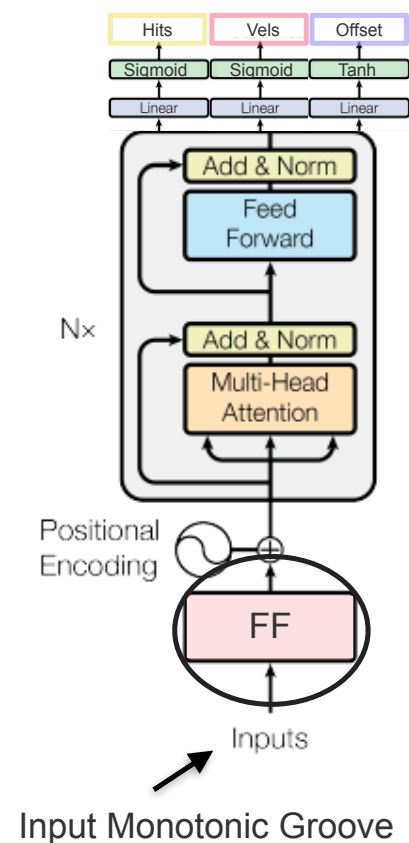
A few limitations:

1. We haven't implemented any control parameters for the generations
2. We can't generate any variations for a given input! (For one input groove, we get a single drum pattern)
3. We can't generate from scratch using this model!

We'll try to make improvements to the architecture by modifying it into a Variational Auto-encoder (VAE)

Drum Pattern

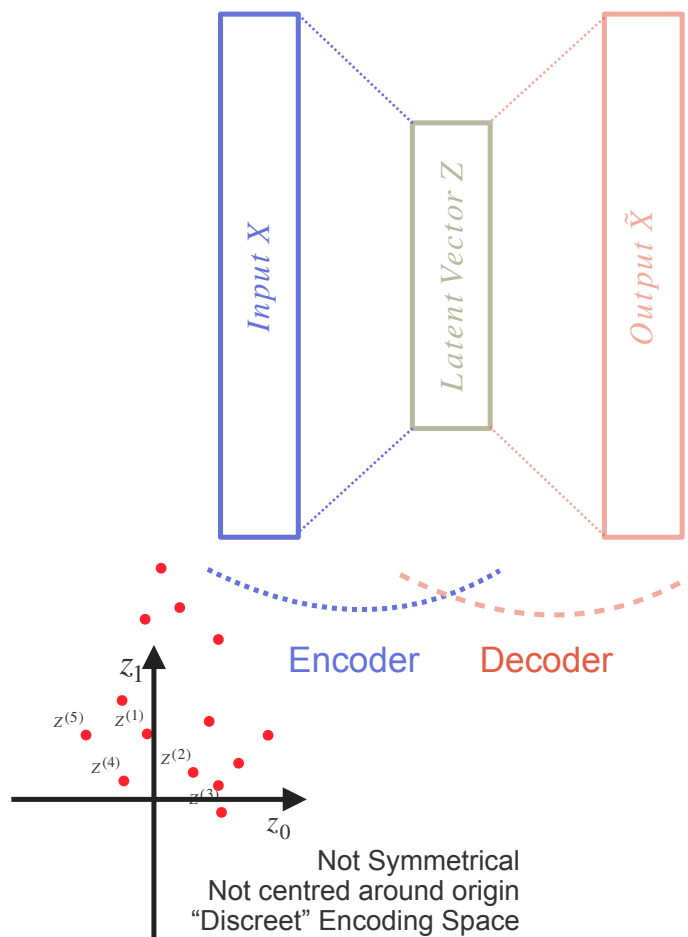
Hi-hat uTiming	0.2	0.42	-0.31	0.42	0.4	0	-0.12	0.18
Snare uTiming							-0.1	
Kick uTiming	0.4		0.23		0.2		-0.35	
Hi-hat Velocity	1	0.7	0.78	0.62	0.73	0.54	0.89	0.51
Snare Velocity							0.64	
Kick Velocity	1		0.4		0.9		0.6	
Hi-hat	1	1	1	1	1	1	1	1
Snare							1	
Kick	1		1		1		1	



0.4		0.23		0.2		-0.35	
1		0.4		0.9		0.6	
1		1		1		1	

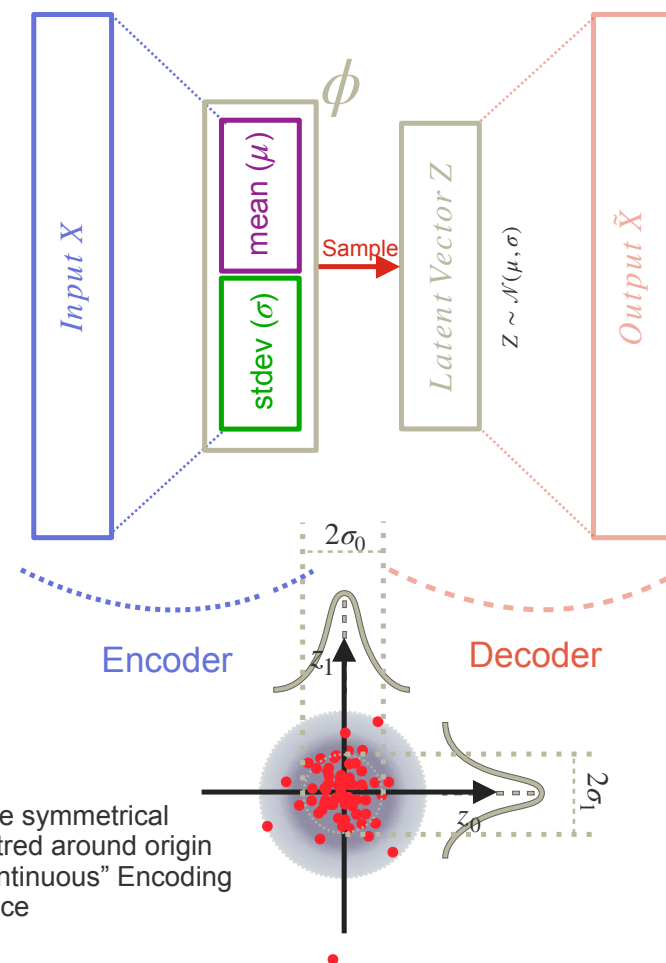
Architecture

Auto-Encoder



$z^{(K)}$: Latent Encoding of K^{th} Input

Variational Auto-Encoder





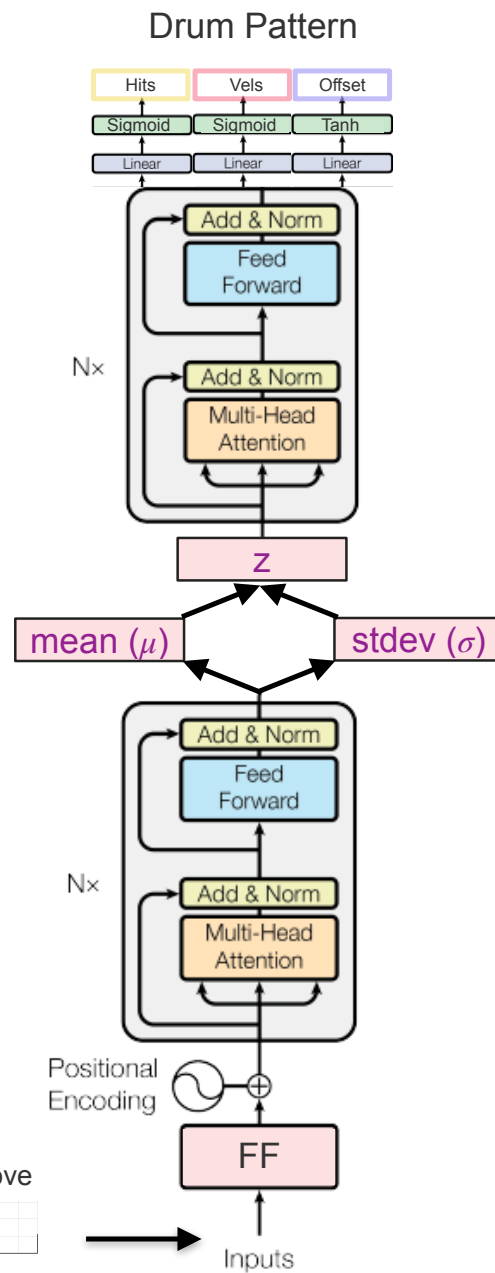
$z^{(K)}$: Latent Encoding of K^{th} Input

Read this if you're interested: <https://towardsdatascience.com/generating-new-faces-with-variational-autoencoders-d13cfcb5f0a8>

Architecture

Input Monotonic Groove

0.4		0.23		0.2		-0.35
1		0.4		0.9		0.6
1		1		1		1



Training

Two sets of Parameters:

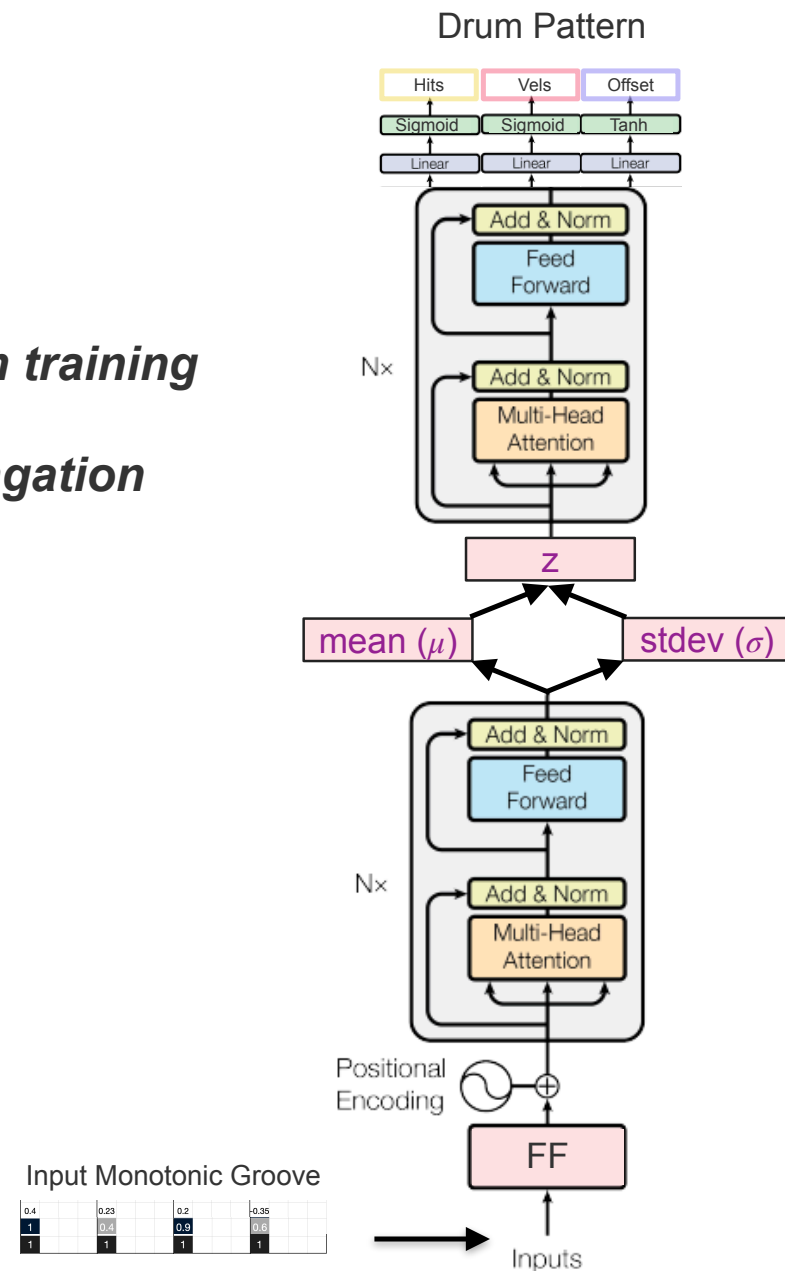
1. Trainable:

1. Weights of Network
2. Biases of Network

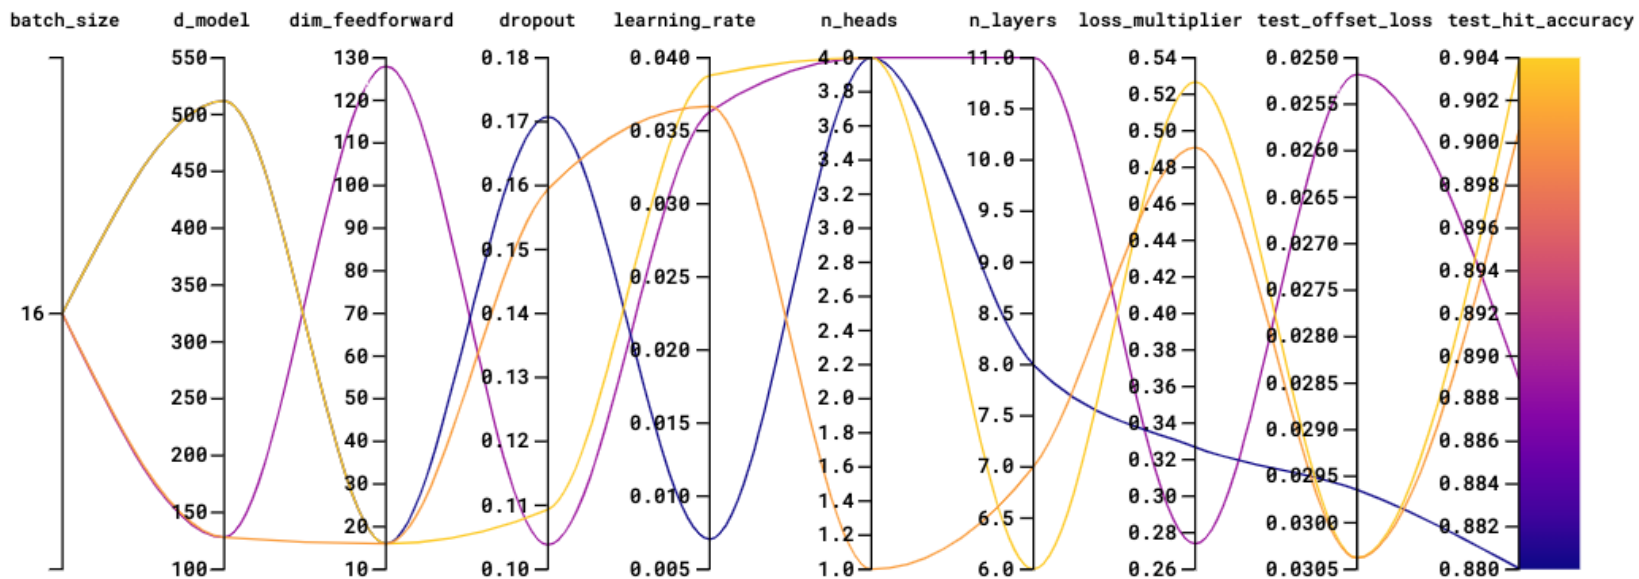
Learnt from training data using back-propagation

2. Hyper-Parameters (Non-trainable):

1. Number of Enc/Dec layers
2. FF dimensionality
3. Z dimensionality
- ...



Training



2. Hyper-Parameters (Non-trainable):

1. Number of Enc/Dec layers
2. FF dimensionality
3. Z dimensionality
- ...

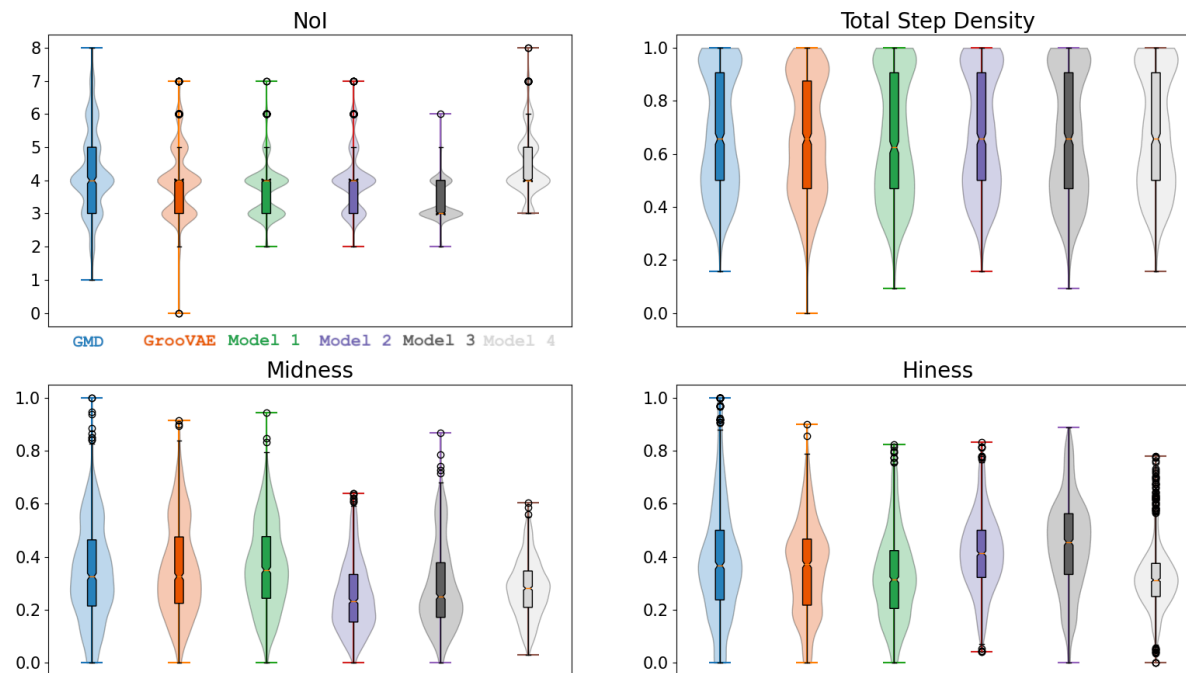
*Tuned by Trial and Error
(AKA Hyper-parameter Tuning)*

Read more:

<https://docs.wandb.ai/guides/sweeps>

Validation by Global Comparison (Feature-based)

Extract musical features and compare models against one another by comparing the global distribution of features

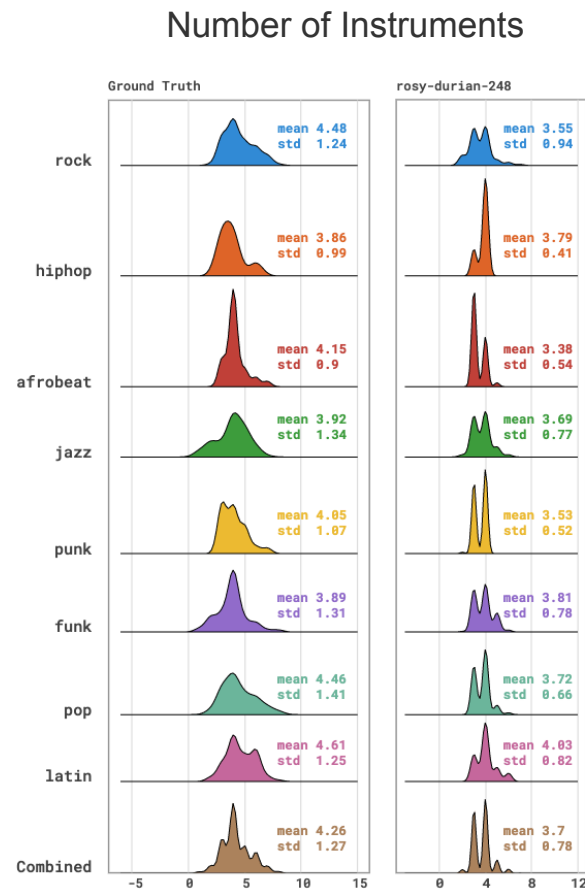


See more here:

<https://wandb.ai/anonmmi/AIMC2022/reports/Absolute-Analysis--VmlldzoxOTU2OTc1>

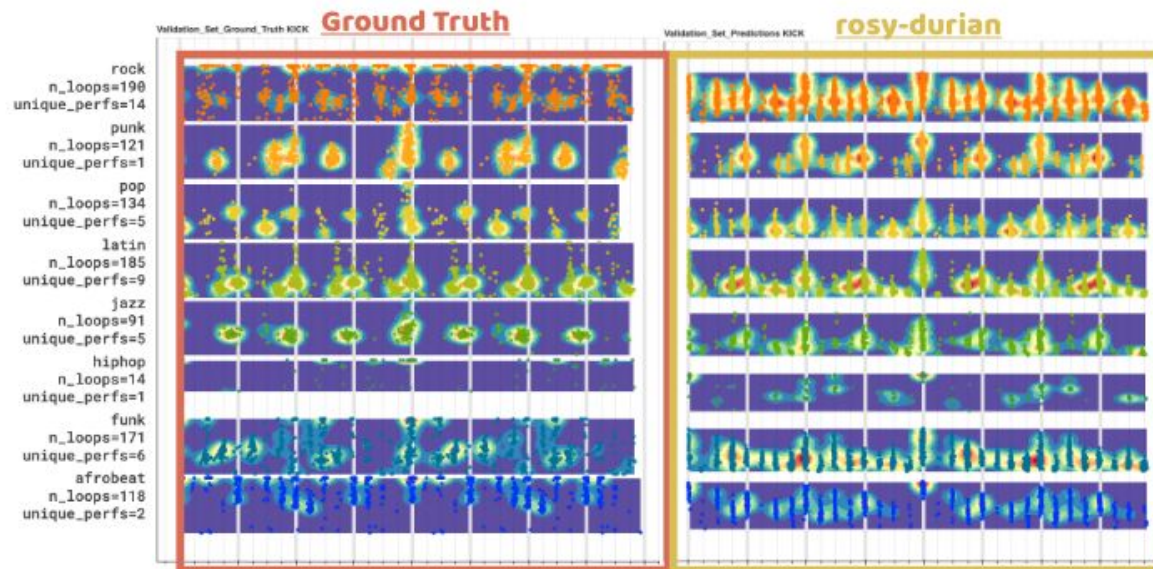
Validation by Global Comparison (Feature-based)

Now do the comparison per genre



Validation by Global Comparison (Velocity Heat-maps)

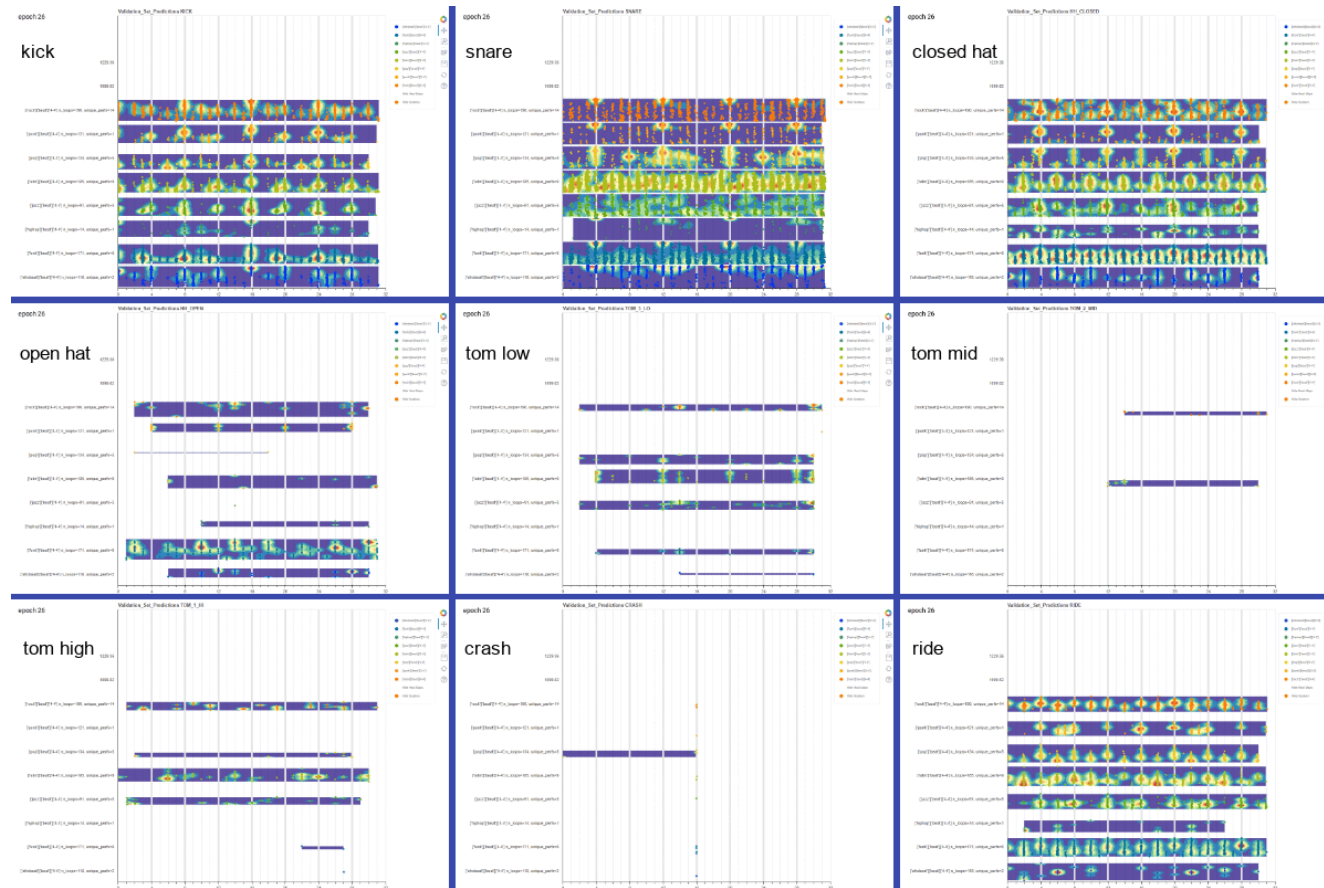
(a) Kick




See more here:

<https://wandb.ai/anonmmi/AIMC2022/reports/Velocity-heatmaps-for-Run-northern-sweep-26---VmIldzoxNTAxNjMy>

Validation by Global Comparison (Velocity Heat-maps)

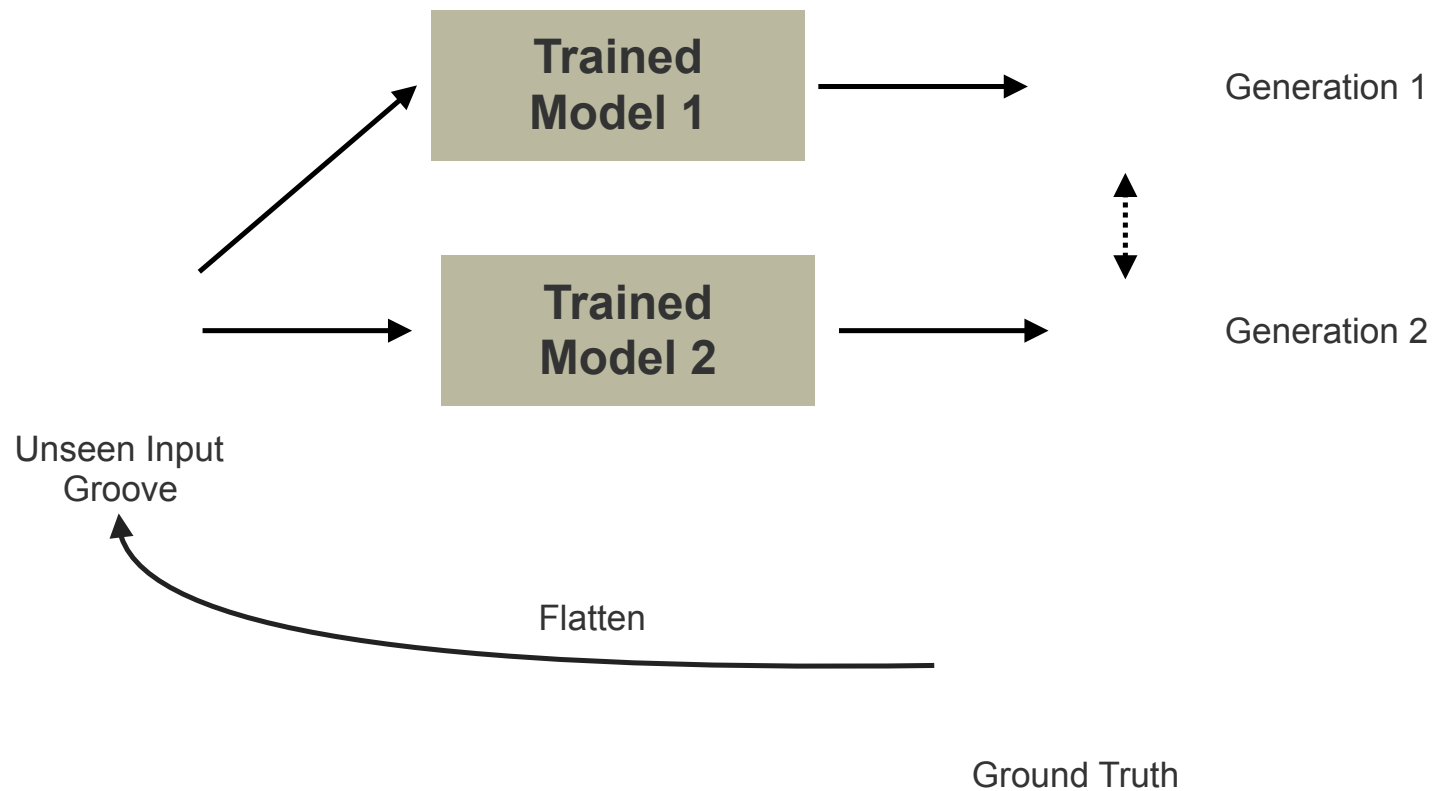


Validation by Global Comparison (Velocity Heat-maps)




Validation by Listening (Subjective)

Use a held-out portion of data to evaluate the performance



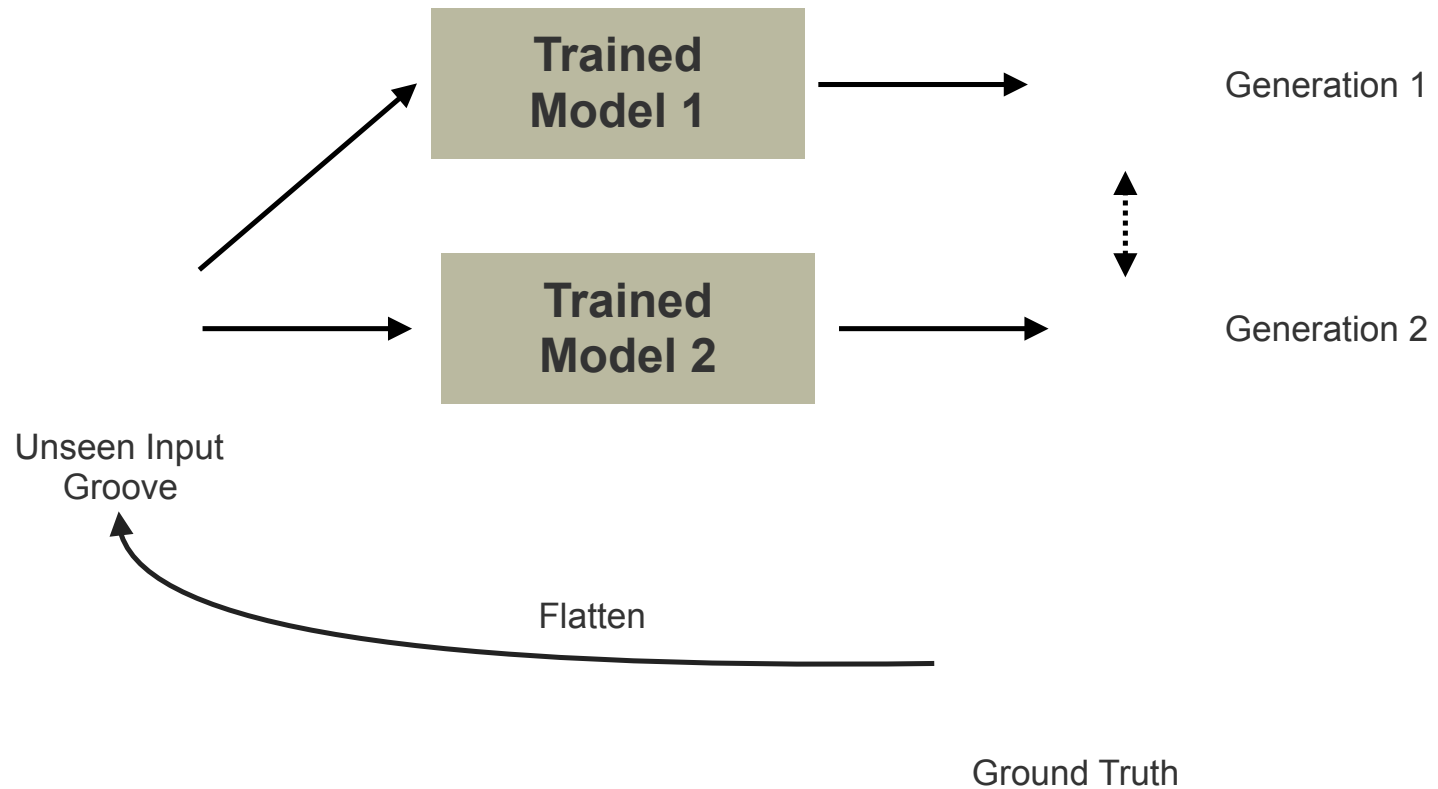
Validation by Listening (Subjective)

Use a held-out portion of data to evaluate the performance




Validation by Listening (Subjective)

Use a held-out portion of data to evaluate the performance




Validation by Listening (Subjective)

Use a held-out portion of data to evaluate the performance




Validation by Listening (Subjective)

Use a held-out portion of data to evaluate the performance



Validation of Embeddings

Use dimension reduction techniques (such as UMAP, T-SNE) to visualize the clustering of embeddings



Check out in supplementary material:

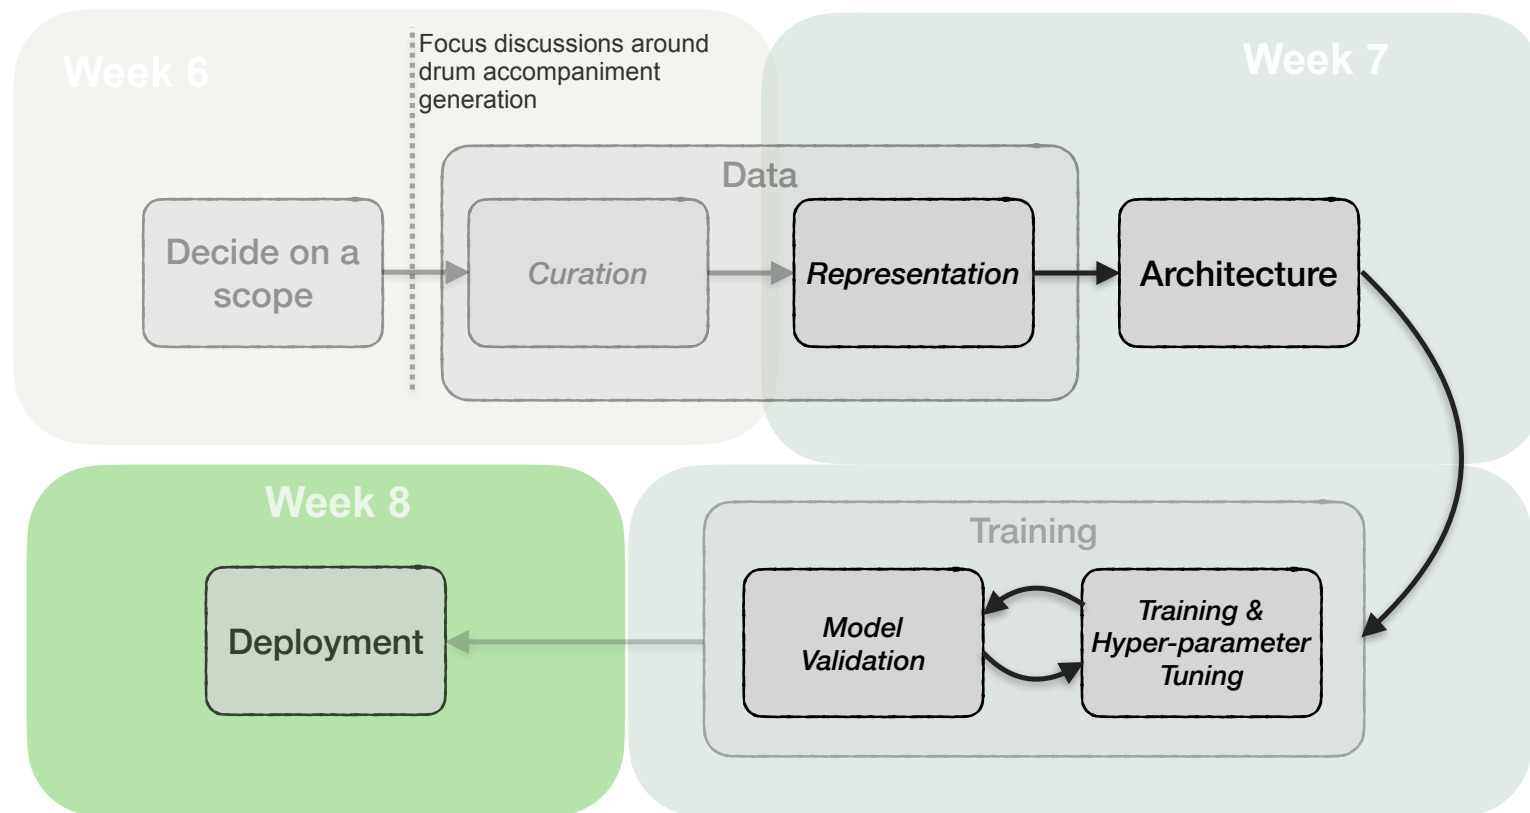
UMAP_drawn_river_6.html
UMAP_noble-field-7.html

Additional Reading

For detailed discussions on topics discussed today:

<https://behzadhaki.com/blog/2022/trainingGrooveTransformer/>

Focus of this week!



RECAP:

Create a real-time drum generation system that accompanies a given instrumental performance

How can we use the trained model in Real-Time?

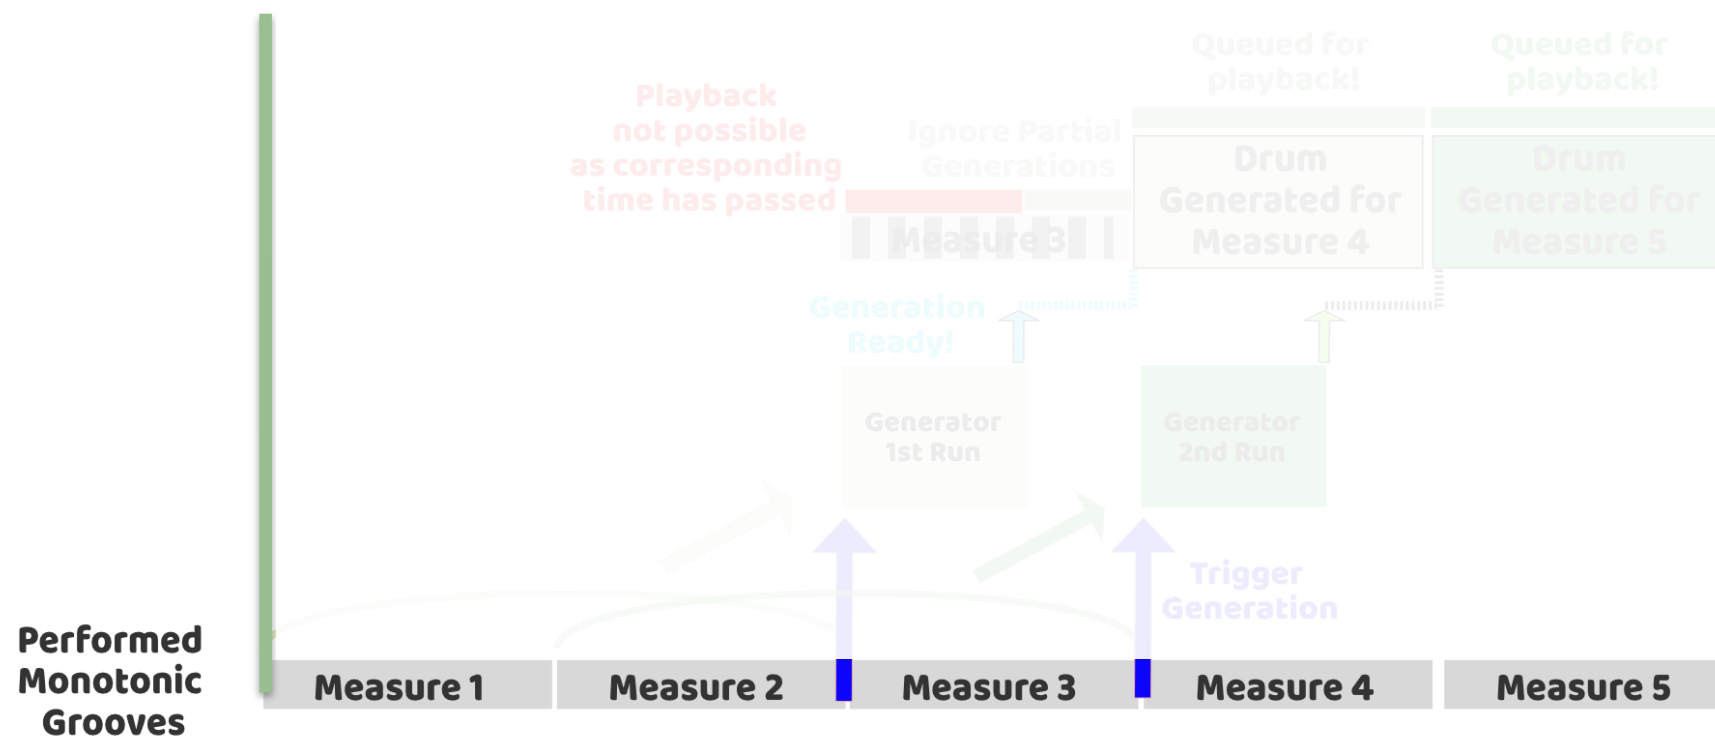
Limitations of the trained model:

- Non-causal Generation (at any time-step, we take into account future events)
- Limited to 2 bars of drums

What if we generate patterns at fixed intervals with a 1-bar delay?

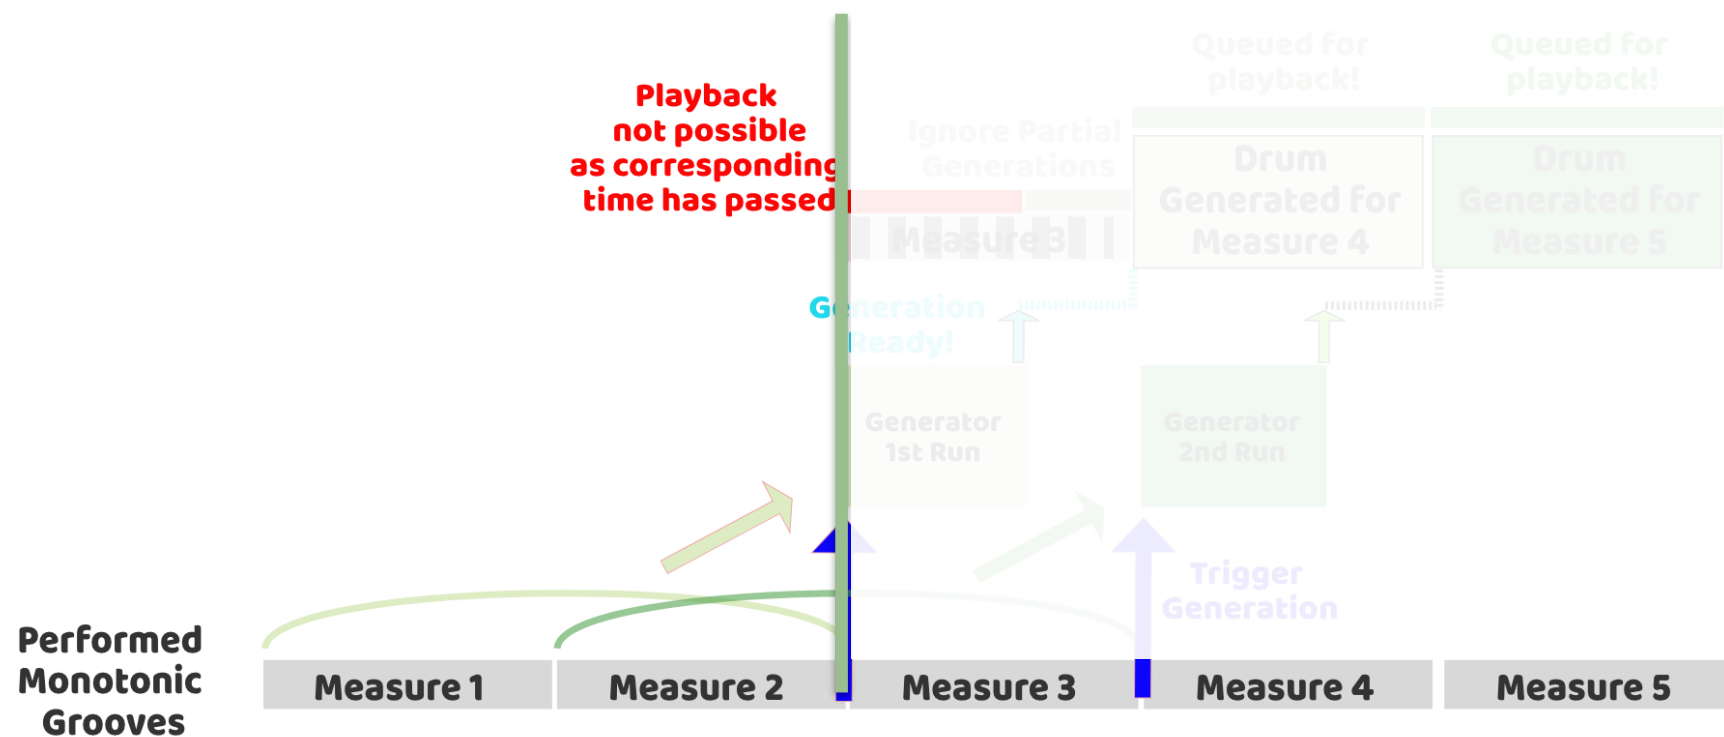
How can we use the trained model in Real-Time?

What if we generate patterns at fixed intervals with a 1-bar delay?



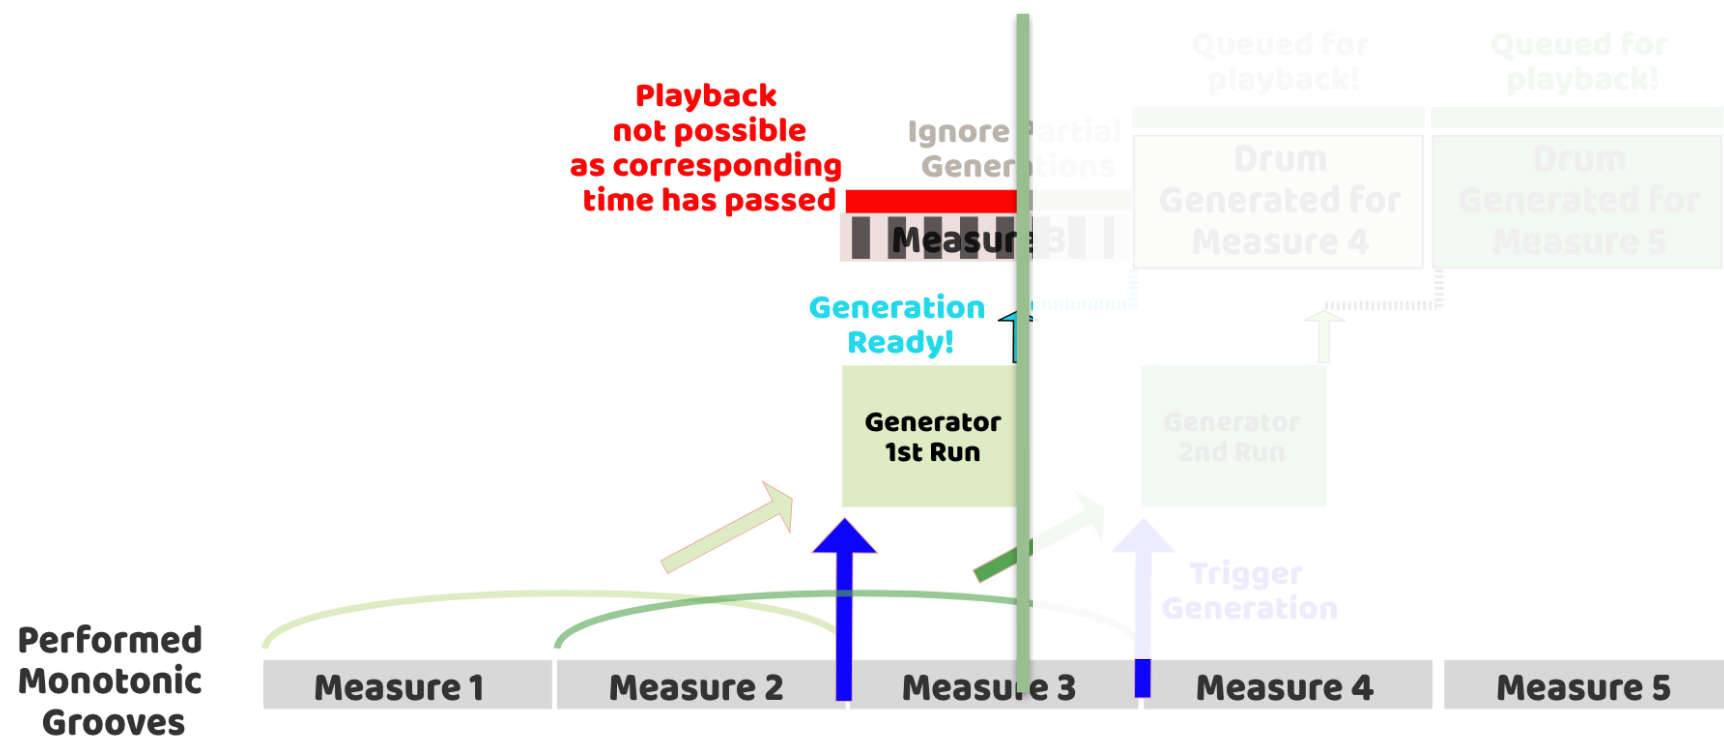
How can we use the trained model in Real-Time?

What if we generate patterns at fixed intervals with a 1-bar delay?



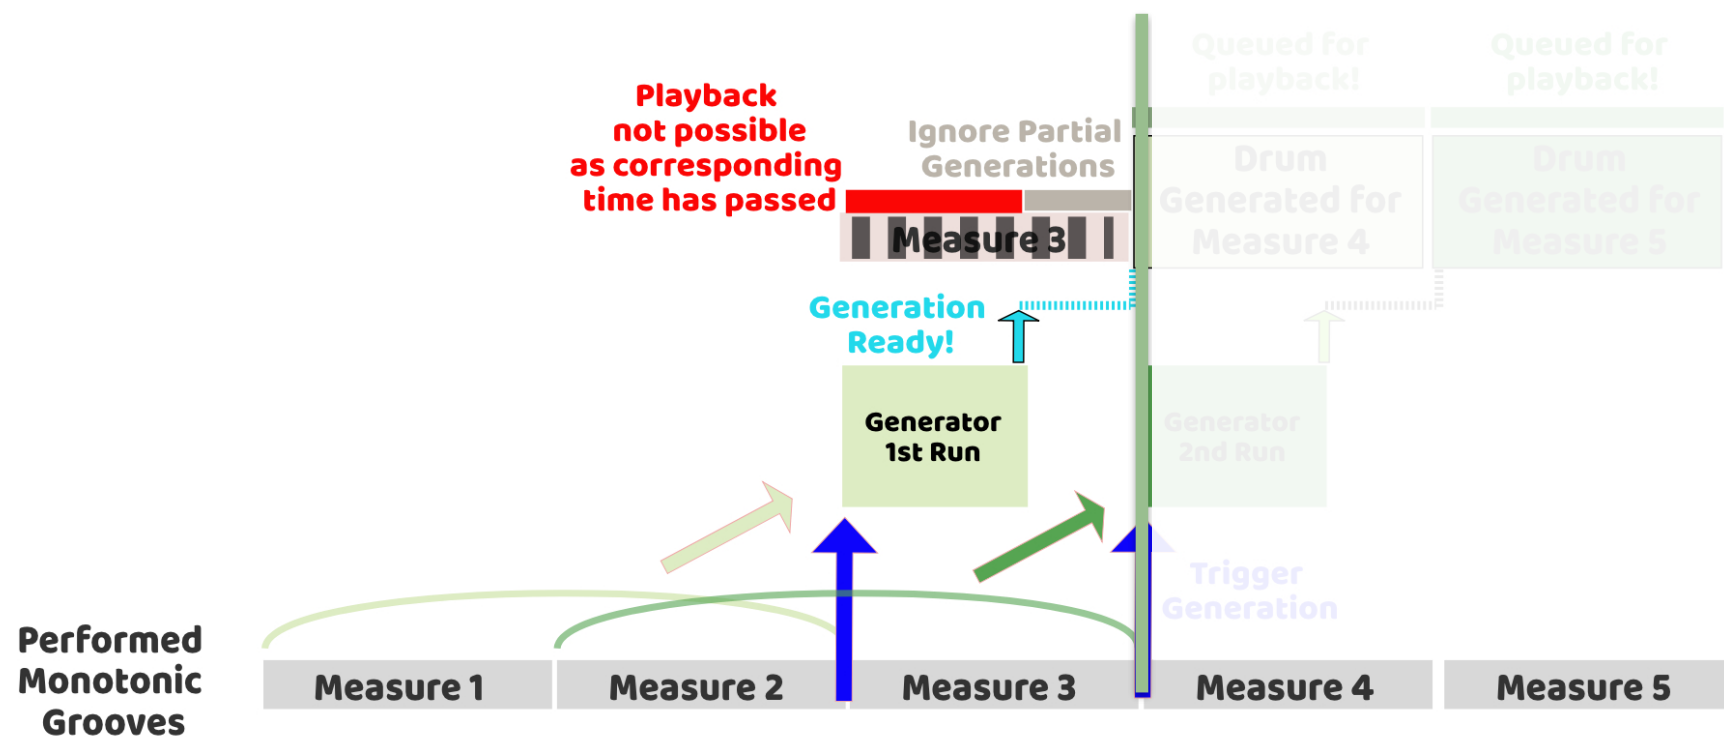
How can we use the trained model in Real-Time?

What if we generate patterns at fixed intervals with a 1-bar delay?



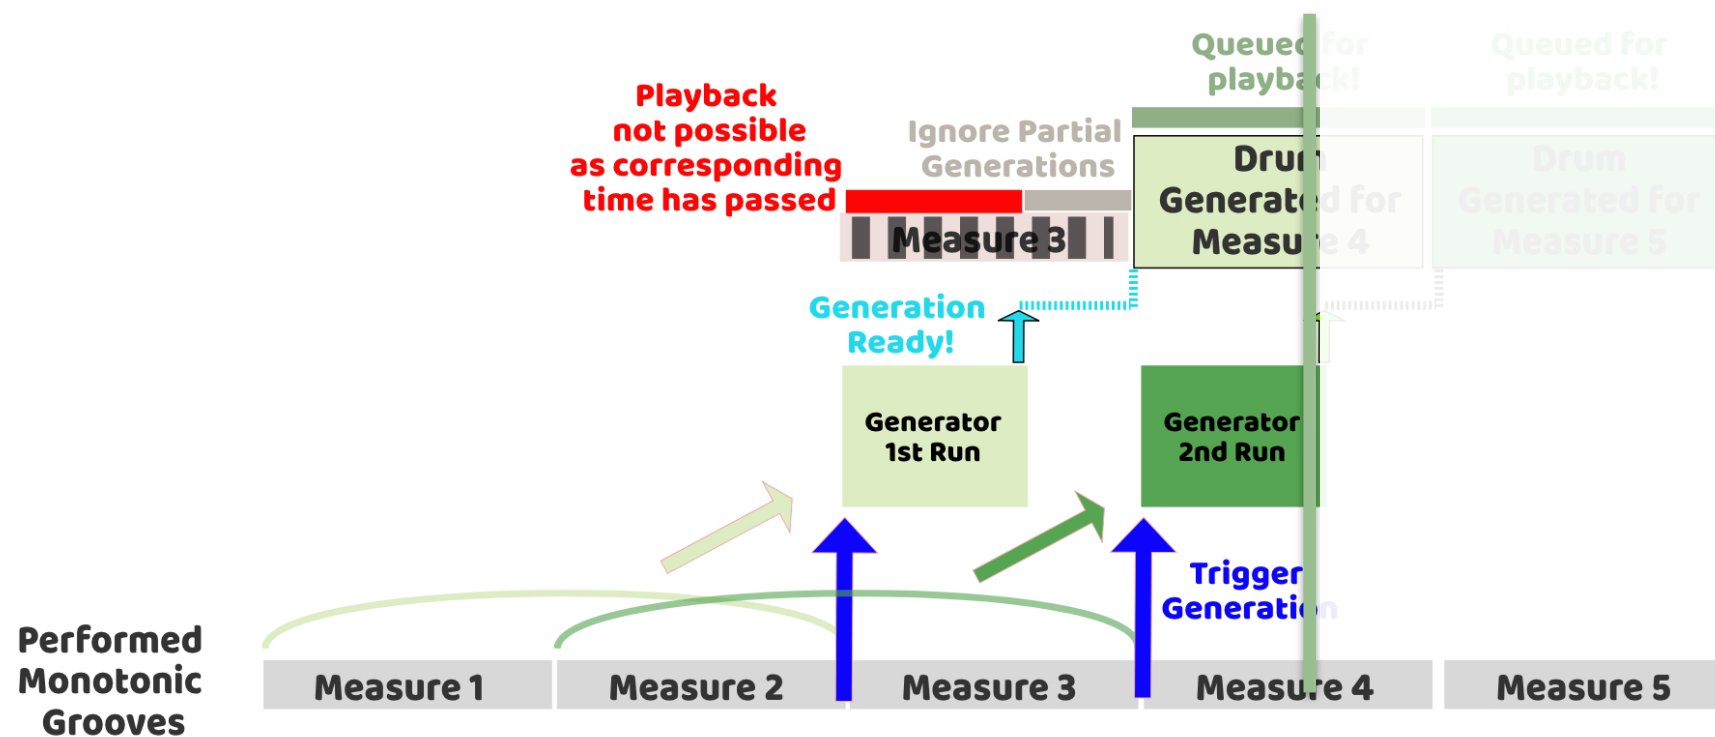
How can we use the trained model in Real-Time?

What if we generate patterns at fixed intervals with a 1-bar delay?



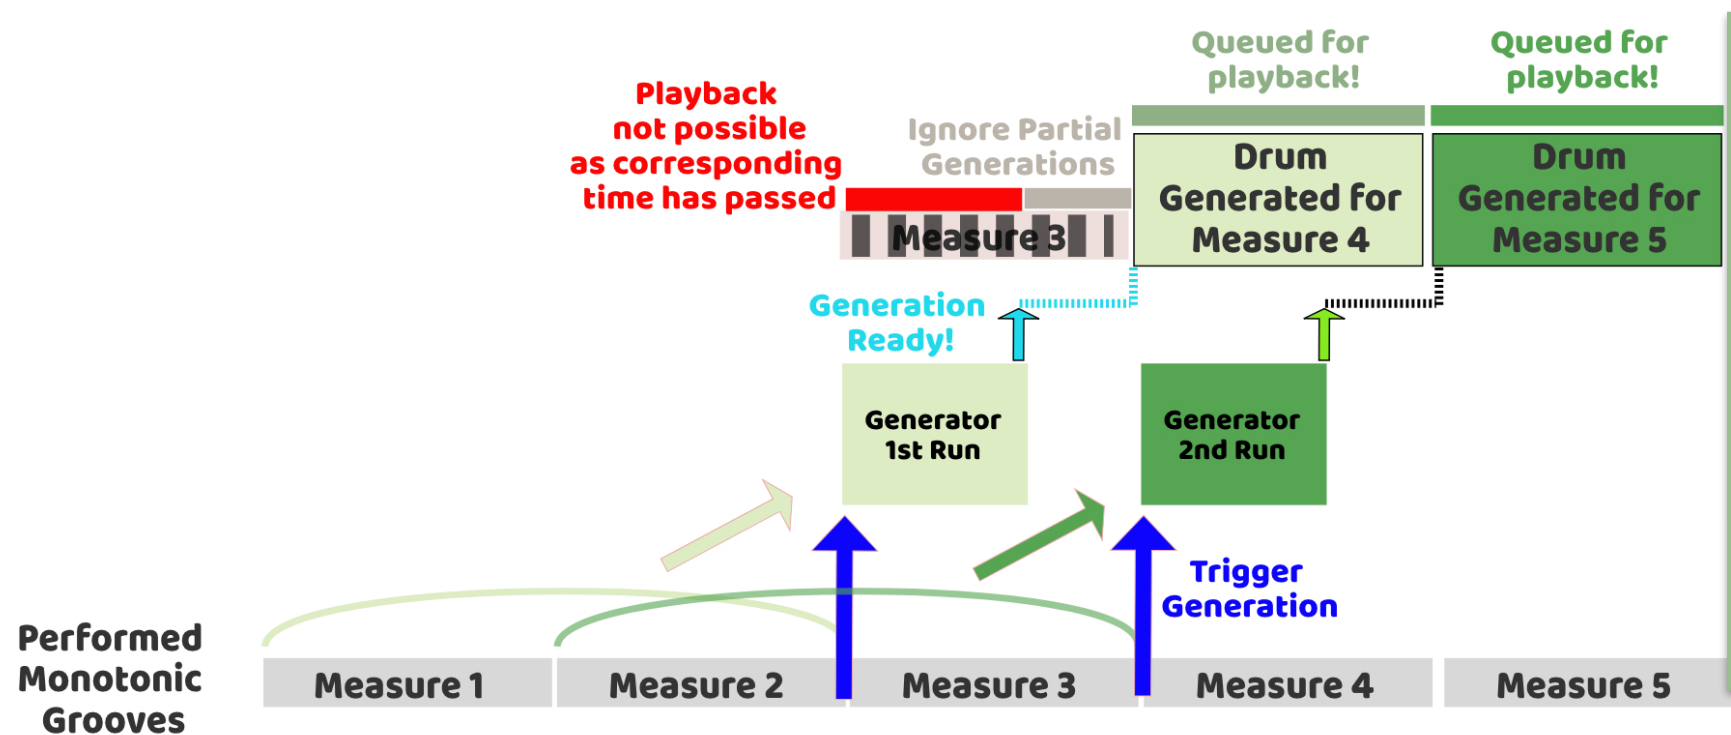
How can we use the trained model in Real-Time?

What if we generate patterns at fixed intervals with a 1-bar delay?

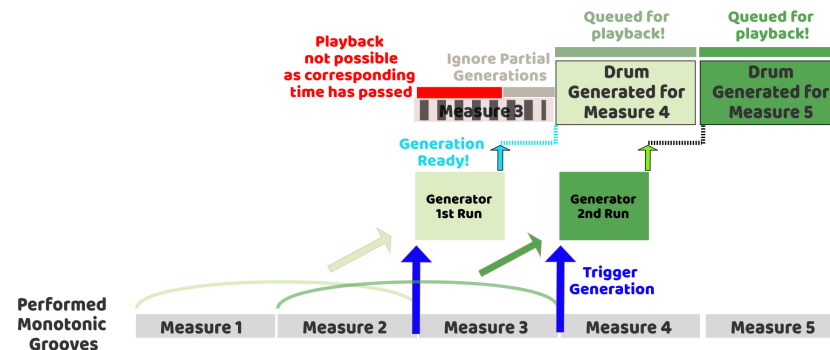


How can we use the trained model in Real-Time?


What if we generate patterns at fixed intervals with a 1-bar delay?



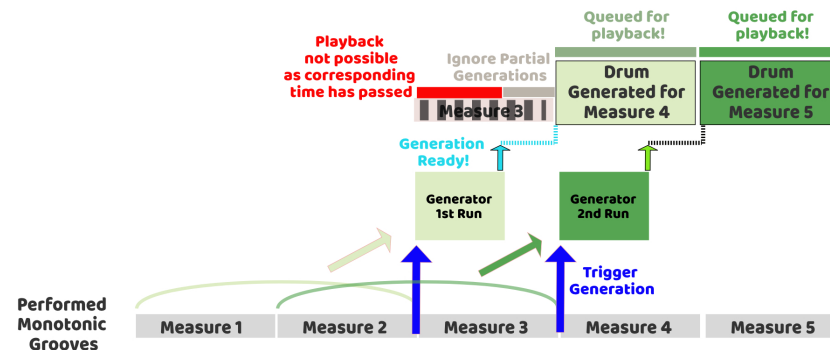
How can we use the trained model in Real-Time?




What if we think of the performance in a looped manner?



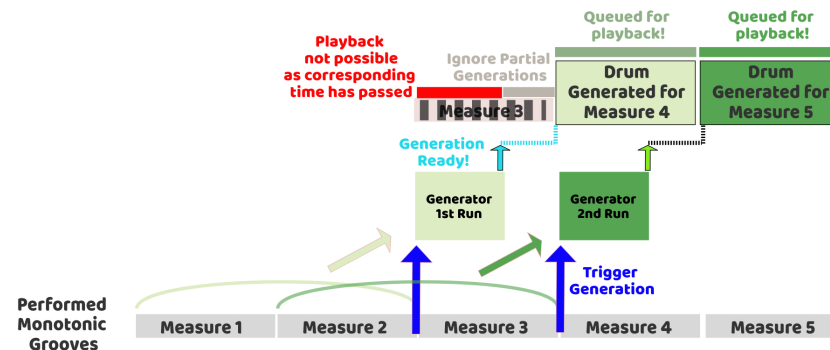
How can we use the trained model in Real-Time?




What if we think of the performance in a looped manner?



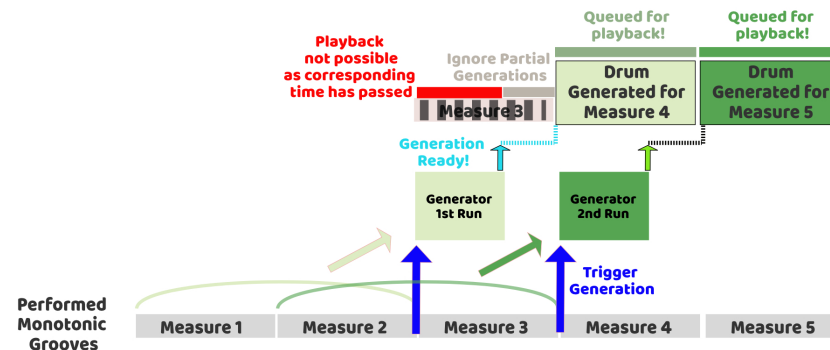
How can we use the trained model in Real-Time?




What if we think of the performance in a looped manner?



How can we use the trained model in Real-Time?



What if we think of the performance in a looped manner?




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

Is inference fast enough for real-time deployment?

- < 10ms (running in python)

Do we need any delay compensation? —> Not in this looping setting




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input




How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input



How can we use the trained model in Real-Time?

One major issue with the model:

Always uses 2-bars of input! —> What about long term coherence of the generations?

Ideally, the system should have been trained on longer sequences

We will try to provide information about the past manually by **OVERDUBBING** the input

